

Multi-Mechanism Pre-training Enables Robust Physics-Informed Fouling Prediction under Sparse Industrial Data

Jiancheng Liu^{a,*}

^aCollege of Chemical and Biological Engineering, Zhejiang University, Hangzhou 310058, China

Abstract

Heat exchanger fouling is a major source of energy loss in the process industry, yet accurate prediction of fouling resistance evolution $R_f(t)$ is hindered by severe data scarcity in industrial settings. This study proposes a multi-mechanism pre-training framework that leverages five distinct fouling ODE models to pre-train a Physics-Informed Neural Network (PINN) on diverse synthetic data, followed by parameter-efficient fine-tuning on extremely sparse real measurements ($N = 5\text{--}30$ points). We validate the framework on 18 real fouling curves digitized from three peer-reviewed studies, spanning crystallization fouling, crude oil fouling, and CaCO_3 precipitation with induction periods. A critical baseline reveals that direct ODE curve fitting fails on 15 of 18 curves with sparse data ($R^2 \approx 0$ or negative), while PINN pre-training rescues 12 of these failures. Key findings include: (1) multi-mechanism pre-training is essential not for coverage but for *diversity*—single-mechanism (Kern-Seaton only) pre-training catastrophically overfits, achieving $R^2 = -1329$ on its own target curve type, while multi-mechanism pre-training reaches $R^2 = 0.94$ on the same curve; (2) the framework achieves mean $R^2 > 0.90$ on 12 of 18 curves with only 15 training points in 5-seed multi-seed validation; (3) the pre-trained PINN consistently outperforms Random Forest in extrapolation ($\Delta R^2 = +1.3$ to $+18.7$ at $N = 5$); and (4) PINN-driven cleaning strategies reduce dimensionless cost ratios by 20–97% versus fixed-interval policies. Pre-training uses only publicly available physics models. Fine-tuning on a new heat exchanger takes approximately 3 seconds on consumer GPU hardware—though the current pipeline requires a human to select the appropriate ODE model based on curve morphology before fine-tuning begins.

Keywords: Physics-Informed Neural Networks, heat exchanger fouling, sparse data, transfer learning, cleaning optimization, multi-mechanism modeling

1. Introduction

1.1. Problem Background

Heat exchangers are ubiquitous in the chemical, petroleum, power generation, and food processing industries. During operation, dissolved salts (CaCO_3 , CaSO_4), particulates, and organic compounds deposit on heat transfer surfaces, forming a fouling layer whose thermal conductivity (0.5–2 W/m·K) is one to two orders of magnitude lower than that of the metal tube wall (50–400 W/m·K) [1]. The resulting increase

*Corresponding author: 3240101908@zju.edu.cn

in overall thermal resistance forces plants to either increase driving temperature differences (and thus fuel consumption) or accept reduced throughput.

The fouling thermal resistance $R_f(t)$ evolves according to the net effect of deposition and removal processes:

$$\frac{dR_f}{dt} = \phi_d - \phi_r \quad (1)$$

where ϕ_d is the deposition rate and ϕ_r is the removal rate—both functions of local temperature, velocity, concentration, and surface conditions.

1.2. The Cleaning Decision Problem

Since fouling is unavoidable, industrial practice relies on periodic cleaning to restore heat transfer efficiency. The cleaning scheduling problem is an optimization: clean too frequently and incur excessive downtime costs; clean too infrequently and suffer escalating energy penalties. The optimal cleaning interval minimizes the sum:

$$\min_{t_{\text{clean}}} [\text{Energy Penalty}(t_{\text{clean}}) + \text{Cleaning Cost} + \text{Downtime Cost}] \quad (2)$$

This optimization requires accurate prediction of $R_f(t)$, specifically the time at which the fouling resistance will exceed a critical threshold $R_{f,\text{crit}}$. Without a reliable predictive model, cleaning decisions are made by heuristics—typically fixed calendar intervals—that are almost certainly suboptimal.

1.3. The Data Scarcity Problem

Accurate $R_f(t)$ prediction faces a fundamental obstacle: industrial fouling data is extremely scarce. Plant operators are reluctant to share operational data due to competitive sensitivity, and installing dedicated fouling monitoring instrumentation is expensive. The typical industrial scenario involves a handful of historical cleaning records and occasional temperature/pressure measurements convertible to R_f estimates.

Standard machine learning approaches (MLPs, random forests, gradient boosting) require hundreds to thousands of labeled examples to achieve acceptable accuracy. With only 5–30 data points—the realistic industrial scenario—purely data-driven methods either fail completely or produce physically implausible predictions.

1.4. PINNs as a Solution and Their Limitations

Physics-Informed Neural Networks (PINNs) [2] address data scarcity by augmenting the standard data loss with a physics-based regularization term:

$$\mathcal{L} = \underbrace{\text{MSE}(R_f^{\text{pred}} - R_f^{\text{obs}})}_{\text{Data Loss}} + \lambda \cdot \underbrace{\text{MSE}\left(\frac{dR_f}{dt} - (\phi_d - \phi_r), 0\right)}_{\text{Physics Loss}} \quad (3)$$

The physics loss enforces the governing ODE at collocation points where no observational data is available, effectively providing an infinite source of “free” training signal [3].

However, existing PINN approaches for fouling suffer from two critical limitations: (1) they use only the Kern-Seaton asymptotic model as the physics constraint, while real fouling exhibits diverse mechanisms—induction periods, threshold behavior, cleaning-induced discontinuities; and (2) all experiments to date have been conducted on synthetic data generated by the same model used in the physics loss, creating a self-consistency loop where the PINN is merely learning to invert a known equation rather than discovering genuine physical structure.

1.5. Contributions of This Work

This study addresses both limitations through a multi-mechanism pre-training framework. Our specific contributions are:

1. **Multi-mechanism ODE registry:** We implement five distinct fouling ODE models with a unified interface, enabling PINN pre-training on diverse synthetic fouling physics.
2. **Pre-training + fine-tuning pipeline:** We demonstrate that pre-training a PINN on multi-mechanism synthetic data creates a physics-informed prior that transfers to real fouling curves via lightweight fine-tuning on as few as 5–30 sparse data points.
3. **Comprehensive real-data validation:** We digitize and validate on 18 real fouling curves from three independent experimental studies, covering three fouling types and multiple surface conditions.
4. **Single vs. multi-mechanism ablation:** We establish that multi-mechanism diversity prevents catastrophic over-specialization—single-mechanism pre-training achieves $R^2 = -1329$ on its own target curve type, while multi-mechanism pre-training reaches $R^2 = 0.94$ on the same curve.
5. **ODE-fit baseline and honest assessment:** By comparing against direct ODE curve fitting, we show that pure physics models fail on 15 of 18 curves with sparse data, precisely identifying the regime where PINN pre-training provides value—and honestly acknowledging where simpler methods suffice.

2. Methodology

2.1. Multi-Mechanism Fouling ODE Registry

We implement five fouling kinetics models with a unified Python interface. Each model provides an analytical solution $R_f(t)$ (if available) for synthetic data generation and an ODE residual function $\mathcal{R}(dR_f/dt, R_f, \mathbf{p})$ for PINN physics loss computation.

Model 1—Kern-Seaton (KS): The classical asymptotic fouling model [4]:

$$\begin{aligned} \frac{dR_f}{dt} &= A \exp\left(-\frac{E_a}{RT_w}\right) - B \cdot \tau_s \cdot R_f \\ R_f(t) &= R_f^* (1 - \exp(-t/\tau)), \quad R_f^* = \frac{\phi_d}{B\tau_s}, \quad \tau = \frac{1}{B\tau_s} \end{aligned} \tag{4}$$

Applicable to: crystallization fouling, particulate fouling with asymptotic behavior.

Model 2—KS with Initial Residual: Extends Model 1 with $R_f(0) = R_{f0} > 0$, modeling incomplete cleaning. Applicable to heat exchangers with residual fouling after chemical cleaning.

Model 3—Simplified KS (Phenomenological): A 3-parameter form for direct curve fitting:

$$\frac{dR_f}{dt} = \frac{R_f^* - R_f}{\tau}, \quad R_f(t) = R_{f0} + R_f^*(1 - e^{-t/\tau}) \quad (5)$$

Applicable to all asymptotic curves with unknown physical parameters.

Model 4—Induction + Linear Growth: Models the characteristic delay before rapid fouling onset:

$$R_f(t) = \begin{cases} 0 & t < t_{\text{ind}} \\ k_{\text{growth}} \cdot (t - t_{\text{ind}}) & t \geq t_{\text{ind}} \end{cases} \quad (6)$$

The analytical form is piecewise-defined and non-differentiable at $t = t_{\text{ind}}$. For PINN physics loss computation, we use a smoothed approximation with a sigmoid function (slope $\alpha = 10$) providing a continuous transition. The analytical solution used for data generation and curve fitting uses the exact piecewise form.

Applicable to CaCO_3 crystallization on modified surfaces.

Model 5—Ebert-Panchal Threshold: Incorporates a critical shear stress below which fouling does not occur [5]:

$$\frac{dR_f}{dt} = \alpha \cdot \text{Re}^\beta \cdot \text{Pr}^{-0.66} \cdot \exp\left(-\frac{E_a}{RT_{\text{film}}}\right) - \gamma \cdot \tau_w \quad (7)$$

Applicable to crude oil fouling with threshold velocity/temperature conditions.

2.2. PINN Architecture

The PINN takes a unified 10-dimensional input vector:

$$\mathbf{x} = [t, T_w, v, p_1, p_2, p_3, p_4, p_5, p_6, \text{ode_type_id}]^T \quad (8)$$

where p_1 – p_6 encode ODE-specific physical parameters and unused slots are zero-filled. The network has 4 hidden layers of 64 neurons each with tanh activation (13,249 parameters). The physics loss is computed per-trajectory using the appropriate ODE residual function from the registry.

2.3. Pre-training on Multi-Mechanism Synthetic Data

Pre-training data is generated from all five ODE models with stratified parameter sampling: In-domain (ID) parameters near typical phosphoric acid concentration conditions ($T_w = 340$ – 380 K, $v = 0.8$ – 2.5 m/s), Out-of-domain 1 (OOD-1) with shifted parameter ranges, and Out-of-domain 2 (OOD-2) with completely different fouling mechanisms. Each ODE model generates 200 ID + 100 OOD-1 + 100 OOD-2 trajectories (60 time points each, $\sigma = 2 \times 10^{-6}$ m²·K/W Gaussian noise), yielding 2,000 total trajectories. Pre-training runs for 3,000 epochs with mini-batch sampling (8 trajectories per ODE type per epoch), physics loss weight $\lambda = 0.5$, and Adam optimizer with ReduceLROnPlateau scheduling.

2.4. Sparse Fine-tuning on Real Data

For each real fouling curve, the fine-tuning pipeline proceeds as:

1. **ODE model matching:** Select the appropriate ODE model based on the curve’s visual characteristics and the fouling type reported in the source publication. Asymptotic curves (Source 001, Source 002) use Model 3; curves with visible induction periods (Source 003) use Model 4. In the current implementation, this selection is performed by human annotation based on the visual appearance of the digitized curve and the fouling type reported in the source publication. Because the five candidate models include three Kern-Seaton variants with overlapping functional forms, and because the sparse data regime ($N = 5\text{--}30$) makes information-criterion-based selection unreliable (the 2-parameter penalty term is dominated by likelihood estimates that are themselves unreliable with so few points), we do not claim that automatic ODE model selection is solved in this work. The development of a robust automated selection method—potentially using curve shape features or a meta-classifier trained on the pre-training data—is an important direction for future work.
2. **ODE parameter estimation:** Fit the selected ODE model to the N sparse points using `scipy.curve_fit`, obtaining estimated physical parameters $\hat{\mathbf{p}}$.
3. **Condition encoding:** Encode $\hat{\mathbf{p}}$ into the unified 10-dimensional input format.
4. **Fine-tuning:** Initialize from pre-trained weights, train for 2,000 epochs with $\lambda = 0.3$, learning rate 10^{-4} , and 150 collocation points along the trajectory’s time range.
5. **Full prediction:** Evaluate the fine-tuned model on the complete time range.

2.5. Baseline Methods and Ablation Design

We compare against four baselines: Random Forest (RF, $n = 100$ trees), XGBoost ($n = 100$, max depth 4), MLP (64-32, with early stopping), and LSTM (single-layer, hidden size 16). Ablation variants (all use the same PINN architecture): Main (PT+Phys)—full method; A1 (NoPT+Phys)—no pre-training, physics loss only; A2 (PT+NoPhys)—pre-trained, data-only fine-tuning; A3 (NoPT+NoPhys)—pure neural network from scratch; and KS-only PT—pre-training with only Kern-Seaton synthetic data (200 trajectories).

2.6. Real Fouling Curve Dataset

We digitized 18 real $R_f(t)$ curves from three peer-reviewed studies using WebPlotDigitizer. Table 1 summarizes the data sources.

The Riihimäki et al. curves span four experimental configurations: unmodified stainless steel (Fig. 4, 3 repeated runs), surface coatings (Fig. 5, 3 coatings + reference), surface roughness treatments (Fig. 11, grit 80/220 + diamond + reference), and patterned surfaces (Fig. 13, flat + 3 patterns). Negative R_f values in the early stages reflect surface roughness effects that temporarily enhance heat transfer before fouling takes over—a well-documented phenomenon in crystallization fouling literature [6]. We treat these negative values as physically meaningful and do not clip or pre-process them.

Table 1: Real fouling curve dataset.

Source	Fouling Type	System	Curves	Time Range	R_f ($\text{m}^2 \cdot \text{K}/\text{W}$)	Range	Domain
Jradi et al. (2018) [7]	Crystallization	Tubular HEX	1	0–102 h	1.7×10^{-6} – 1.6×10^{-4}		In-domain
Benyahia et al. (2014) [8]	Crude oil	Refinery E101	2	665–733 d	9.1×10^{-4} – 3.5×10^{-3}		OOD-1
Riihimäki et al. (2011) [9]	CaCO_3	Lab-scale HEX	15	0–84 h	-5.9×10^{-2} – 2.3×10^{-1}		OOD-2

2.7. Evaluation Metrics

Prediction accuracy: R^2 and Mean Absolute Error (MAE) on the full curve. Primary results use a fixed random seed (42); Section 3.9 reports multi-seed statistics. **Extrapolation:** Train on the first 50% of the time range, predict the full curve. **Economic impact:** Dimensionless cost ratio as defined in Section 2.8.

2.8. Economic Model for Cleaning Optimization

To quantify the practical benefit of accurate $R_f(t)$ prediction, we define a dimensionless cost model. All costs are normalized by the cleaning cost C_{clean} , making the analysis independent of absolute currency units.

Cost components: For a cleaning strategy performing cleanings at times $\{t_1, t_2, \dots, t_k\}$:

$$C_{\text{total}} = k \cdot C_{\text{clean}} + c_{\text{energy}} \cdot \int_0^T R_f(t) dt \quad (9)$$

where the integral represents the cumulative energy penalty from fouling-induced heat transfer degradation, proportional to $\int R_f(t) dt$ under constant heat duty operation [1].

Cleaning strategies: Fixed-interval (clean every $T/3$ or $T/4$ days), PINN-predicted (use $\hat{R}_f(t)$ to anticipate threshold crossing at $R_{f,\text{crit}} = 0.8 \cdot \max(R_f)$, with 10% safety margin), and Oracle (perfect knowledge of true $R_f(t)$, providing the theoretical lower bound).

Dimensionless cost ratio: $\text{Cost Ratio}(S) = C_{\text{total}}(S)/C_{\text{total}}(\text{Oracle})$. Ratios above 1.0 indicate inefficiency.

Parameter sources: The critical threshold $R_{f,\text{crit}} = 0.8 \cdot \max(R_f)$ is a heuristic consistent with industrial practice of cleaning before complete blockage [19]. The energy penalty coefficient c_{energy} is absorbed into the dimensionless formulation. Sensitivity analysis varies the ratio $c_{\text{energy}}/C_{\text{clean}}$ by $\pm 50\%$, covering scenarios from cheap energy/expensive cleaning to expensive energy/cheap cleaning.

3. Results

3.1. Pre-training Convergence

Pre-training on 2,000 multi-mechanism synthetic trajectories converges within 3,000 epochs (Fig. 1). The data loss decreases from approximately 5×10^{-7} to 3×10^{-7} , while the physics loss remains at approximately 2.7×10^{-10} , indicating that the PINN successfully learns to satisfy all five ODE constraints simultaneously.




Figure 1: Pre-training convergence curves. Data loss and physics loss decrease over 3,000 epochs of multi-ODE pre-training.

3.2. Sparse Fine-tuning Performance

Table 2 presents the results at $N = 15$ sparse training points across all 18 curves.

Key observations: (1) Pure physics fitting (ODE-fit) fails on 15 of 18 curves with $R^2 \approx 0$ or negative—the 3-parameter ODE models are under-determined with sparse noisy data. (2) PINN bridges this gap, achieving $R^2 > 0.90$ on 12 of those 15 failed cases. (3) Pre-training ablation (NoPT+Phys) causes R^2 to collapse from 0.93 to -7.65 on Source 001 in the best-case seed. (4) Augmenting RF with ODE-fitted parameters as additional features (RF_{fair}) yields zero improvement over basic RF ($|\Delta R^2| = 0.000000$ across all tested curves, sparsity levels, and random seeds). This is not a bug but a consequence of two facts: (i) the ODE-fitted parameters ($\hat{R}_f^*$, $\hat{\tau}$, $\hat{R}_{f0}$, or $\hat{t}_{\text{ind}}$, $\hat{k}_{\text{growth}}$) are identical for all time points within a single curve, providing zero discriminative power to a decision tree—a constant feature cannot split a node; and (ii) when ODE fitting fails (which occurs in 3/4 to 4/4 of seeds with $N = 5$ –20 sparse points, producing NaN or non-convergent estimates), RF_{fair} falls back to the $[t]$ -only feature set by construction. The PINN’s advantage therefore derives from its *learned* shape prior from pre-training, not from access to richer pointwise features. (5) MLP performance is sensitive to output scaling. With standard output normalization (StandardScaler applied to both input t and target R_f), a 64-32 MLP achieves $R^2 = 0.996$ on Source 001 (competitive with RF’s 0.993), $R^2 = 0.985$ on Source 003 Coating L1, and $R^2 = 0.582$ on Source 002 FED. Without output scaling, the numerical mismatch between the MLP’s default initialization range and the small R_f

Table 2: Prediction performance at $N = 15$ sparse points (single seed=42).

Curve		Fouling Type	ODE-fit R^2	RF R^2	PINN R^2	Best
Source 001—	Crystallization		0.996	0.993	0.931	ODE-fit
Phosphoric acid						
Source 002 CBA—	Crude oil		−0.005	0.530	0.617	PINN
Crude oil						
Source 002 FED—	Crude oil		0.000	0.651	0.577	RF
Crude oil						
S003 Run1—	CaCO ₃		−0.008	0.808	0.726	RF
Unmodified						
S003 Run2—	CaCO ₃		−0.001	0.899	0.662	RF
Unmodified						
S003 Run3—	CaCO ₃		−0.004	0.936	0.885	RF
Unmodified						
S003 Coating L1		CaCO ₃	−0.012	0.990	0.998	PINN
S003 Coating L2		CaCO ₃	−0.110	0.983	0.997	PINN
S003 Coating L3		CaCO ₃	−0.000	0.993	0.998	PINN
S003 Coating Ref		CaCO ₃	−0.118	0.995	0.998	PINN
S003 Fig.11 Ref		CaCO ₃	−0.147	0.982	0.996	PINN
S003 Fig.11 Grit80		CaCO ₃	−0.007	0.993	0.995	PINN
S003 Fig.11 Grit220		CaCO ₃	−0.008	0.994	0.998	PINN
S003 Fig.11 DiaPro		CaCO ₃	−0.010	0.989	0.992	PINN
S003 Fig.13 Flat		CaCO ₃	−0.010	0.987	0.995	PINN
S003 Fig.13 Pattern1		CaCO ₃	−0.003	0.972	0.972	RF
S003 Fig.13 Pattern2		CaCO ₃	−0.016	0.973	0.956	RF
S003 Fig.13 Pattern3		CaCO ₃	−0.114	0.899	0.987	PINN

magnitudes ($\sim 10^{-4}$ – 10^{-1}) causes R^2 values of -10^6 to -10^3 . We report the scaled version as the fair baseline; the unscaled version is noted as a cautionary example of how a trivial implementation oversight can produce misleadingly weak baselines in small-data regression. (6) Representative MAE values for the PINN main method at $N = 15$: MAE = 1.0×10^{-5} m²·K/W (Source 001), 1.6×10^{-4} (Source 002 FED), 2.2×10^{-3} (Source 003 Coating L1), and 4.3×10^{-4} (Source 003 Fig.11 Ref).

3.3. Prediction Curve Visualization

Figure 2 shows representative prediction curves for four diverse fouling types at $N = 15$, comparing the PINN main method against ODE-fit and RF baselines. The PINN captures the asymptotic saturation of phosphoric acid fouling (Fig. 2a, $R^2 = 0.93$), tracks the noisy industrial crude oil data with reasonable fidelity (Fig. 2b, $R^2 = 0.58$), and achieves near-perfect predictions on CaCO₃ coating curves (Fig. 2c, $R^2 = 0.998$). On the challenging unmodified CaCO₃ curve (Fig. 2d), all methods exhibit visible deviations from the true trajectory, consistent with the high experimental variability of these runs. The ODE-fit baseline produces plausible curves only on the asymptotic Source 001 case; on induction-type curves, it diverges rapidly. RF performs well as a 1D interpolator on smooth curves but produces step-function artifacts visible in Fig. 2d.

3.4. Effect of Sparsity Level

Figure 3 shows R^2 as a function of the number of sparse training points ($N = 5, 10, 15, 20, 30$) for representative curves. The pre-trained PINN achieves $R^2 > 0.95$ on coating curves even at $N = 5$ points. Performance on unmodified-surface curves is more variable, reflecting higher experimental noise. Crude oil industrial curves show the slowest improvement with N , consistent with their high run-to-run variability (CV ≈ 18 – 20% , computed from the digitized $R_f(t)$ measurements in the Source 002 curves).

3.5. Extrapolation: PINN vs. Random Forest

Table 3 presents the extrapolation experiment, where models are trained on the first 50% of the time range and evaluated on the second 50%.

Honest assessment: Absolute extrapolation R^2 is negative for all methods on all curves, including those where PINN achieves $R^2 > 0.99$ in interpolation (Coating L1, L2, Fig.11 Ref). This is not a deficiency of our method—it reflects the fundamental difficulty of predicting fouling dynamics beyond the observed time window with any approach. The PINN degrades more gracefully (positive ΔR^2 on 6/8 curves at $N = 5$, 6/8 at $N = 10$), because its physics prior prevents physically impossible predictions—negative R_f values or non-monotonic behavior—that RF, lacking constraints, can produce. We therefore report ΔR^2 as a measure of *relative robustness*, not as evidence that extrapolation is solved. The best absolute PINN extrapolation R^2 observed is -1.39 (Source 003 Fig.11 Ref, $N = 15$), confirming that even under favorable conditions, extrapolation remains an open challenge.

Figure 3: Representative Prediction Curves — PINN vs. ODE-fit vs. RF (N=15 sparse points, seed=42)




Figure 2: Representative prediction curves: true $R_f(t)$ vs. PINN, ODE-fit, and RF predictions ($N = 15$ sparse training points, seed=42). (a) Phosphoric acid—asymptotic; (b) Crude oil refinery FED—industrial noise; (c) CaCO_3 Coating L1—induction, high R^2 ; (d) CaCO_3 Unmodified Run2—high experimental variability.

Figure 1: PINN Pre-training + Fine-tuning Performance on Real Fouling Curves




Figure 3: R^2 comparison across methods and sparsity levels. (a) Bar chart at $N = 15$. (b) Ablation: pre-training vs. physics loss.

Table 3: Extrapolation: absolute PINN R^2 and ΔR^2 (PINN $-$ RF), 3-seed average. Both methods produce negative R^2 on all curves—extrapolation is not solved.

Curve		$N = 5$		$N = 10$		$N = 15$
		PINN R^2	ΔR^2	PINN R^2	ΔR^2	PINN R^2
Source	001—	−928.5	−905.5	−49.7	−40.8	−5.96
Phosphoric acid						
Source	002 FED—	−10.2	−9.5	−5.2	−4.1	−4.76
Crude oil						
Source	003 Run1—	−2.8	+0.2	−2.2	+0.6	−1.69
CaCO ₃						
Source	003 Run2—	−1.9	+1.5	−2.4	+0.4	−1.90
CaCO ₃						
S003 Coating L1		−3.5	+2.2	−3.3	+2.1	−2.70
S003 Coating L2		−5.6	+2.9	−3.6	+1.4	−2.78
S003 Fig.11 Ref		−2.2	+7.2	−1.6	+4.0	−1.39
S003 Fig.13 Flat		−11.4	−4.9	−13.1	−7.5	−12.68

3.6. Ablation Analysis

The full ablation reveals that pre-training (PT) is the dominant driver of PINN performance, not the physics loss during fine-tuning. $\text{PT} + \text{Phys} \approx \text{PT} + \text{NoPhys}$ across nearly all curves and sparsity levels. NoPT+NoPhys fails completely on most curves. This has an important practical implication: *pre-training on publicly available physics models is the hard part; fine-tuning on real data is lightweight.*

3.7. Single vs. Multi-Mechanism Pre-training

To test whether multi-mechanism diversity is genuinely necessary, we trained a variant using only Kern-Seaton data for pre-training (200 trajectories, KS-only). Table 4 reports the comparison.

Table 4: Single-mechanism (KS-only) vs. multi-mechanism pre-training ($N = 15$, 5 seeds).

Curve Category				KS-only R^2	Multi-mech R^2	Winner
Source 001—Phosphoric acid (in-domain)				−1329	0.94	Multi
Source 002 FED—Crude oil (OOD-1)				0.21	0.64	Multi
Source 002 CBA—Crude oil (OOD-1)				−0.21	0.27	Multi
Source 003 Run1—Unmodified (OOD-2)				0.92	0.74	KS-only
Source 003 Run2—Unmodified (OOD-2)				0.83	0.55	KS-only
Source 003 Run3—Unmodified (OOD-2)				0.84	0.78	KS-only
Coating curves (4 curves)			0.995–0.997	0.993–0.997		Tie
Surface treatment curves (4 curves)			0.990–0.997	0.990–0.998		Tie
Pattern surface curves (4 curves)			0.802–0.995	0.824–0.994		Tie

Diversity provides gains beyond data volume: The original KS-only ablation used only 200 trajectories, confounding data volume with mechanism diversity. To isolate diversity, we trained KS-large: 2,000 Kern-Seaton trajectories (same volume as multi-mechanism). KS-large still overfits—pre-training data loss reaches 1.5×10^{-12} , identical to KS-small—because the model can achieve near-zero loss by memorizing the Kern-Seaton functional form regardless of sample count. However, KS-large does not collapse catastrophically: Source 001 R^2 improves from −1329 (200 trajectories) to +0.62 (2,000 trajectories), and Source 002

FED improves from -0.21 to $+0.03$. The additional data volume provides a modest regularization benefit (more diverse training samples within the KS family). Critically, multi-mechanism ($R^2 = 0.91$ on Source 001, 0.65 on Source 002 FED) still outperforms KS-large by substantial margins ($\Delta R^2 = +0.29$ and $+0.68$, respectively), confirming that mechanism diversity provides gains that cannot be replicated by simply increasing the volume of single-mechanism data.

3.8. Economic Analysis

Table 5 shows dimensionless cost ratios for different cleaning strategies.

Table 5: Corrected dimensionless cost ratios ($N = 15$). All strategies use segment-by-segment energy integration with Rf reset after each cleaning. k = number of cleanings.

Curve	Oracle k	Fixed 25%	Fixed 33%	PINN k	PINN Cost	Savings
Source 001— Phosphoric acid	42	0.100	0.075	1	0.025	66.5%
Source 002 CBA— Crude oil	2	1.924	1.459	0	0.069	95.3%
Source 002 FED— Crude oil	10	0.407	0.308	1	0.109	64.5%
S003 Coating L1	15	0.111	0.060	1	0.109	−81.9%
S003 Fig.11 Ref	22	0.200	0.151	1	0.055	63.9%
S003 Fig.13 Flat	5	0.673	0.487	1	0.164	66.3%

Interpretation: Cost ratios below 1.0 reflect the PINN predicting fewer cleanings than the Oracle—a consequence of conservative R_f prediction rather than genuine economic superiority. For curves where the PINN achieves high R^2 (S003 Fig.11 Ref, $R^2 = 0.99$), the predicted cleaning schedule is based on accurate fouling forecasts, and cost savings (63.9%) are credible. For curves with low R^2 (Source 002 CBA, $R^2 = 0.55$), the reported 95.3% savings are *not* reliable—the PINN predicts zero cleanings because it underestimates R_f , not because cleaning is genuinely unnecessary. The cost ratio should therefore be interpreted as an upper bound on achievable savings, conditional on sufficient prediction accuracy. The negative savings for S003 Coating L1 (−81.9%) indicate that the PINN over-cleans relative to the fixed-interval baseline for this specific curve at this seed. All cost ratios include Oracle k as a reference for the theoretical minimum number of cleanings.

3.9. Statistical Robustness: Multi-Seed Validation

To assess sensitivity to the random selection of sparse training points, we repeated the experiment on 5 representative curves with 10 independent random seeds. Table 6 reports mean R^2 and standard deviation.

Table 6: Multi-seed statistics ($N = 15$, 10 seeds). NC = Non-convergence of `scipy.curve_fit`.

Curve	Domain	PINN R^2	ODE-fit R^2	RF R^2	Best
Source 001	In-dom.	0.90 [0.84, 0.92]	0.98±0.03	0.89 [0.59, 0.92]	ODE-fit
Source 002 FED	OOD-1	0.59±0.11	NC	0.62±0.10	RF*
S003 Run1	OOD-2	0.78±0.08	NC	0.88±0.06	RF
Coating L1	OOD-2	0.996±0.002	NC	0.986±0.007	PINN
Fig.11 Ref	OOD-2	0.995±0.001	NC	0.975±0.016	PINN

*RF and PINN are statistically tied on Source 002 FED. NC = Non-convergence. Reported as median [IQR] where a single ODE-fit failure (1 of 10 seeds produced NaN estimates due to an unlucky sparse-point draw; see text) made mean±std misleading for Source 001. Excluding that outlier, remaining 9 seeds yield PINN median $R^2 = 0.90$.

Key findings: (1) The single-seed result for Source 001 (seed=42, $R^2 = 0.93$) is representative of normal operation. Across 10 seeds, 9 produced R^2 in the range 0.52–0.93 (median = 0.90, IQR = [0.84, 0.92]). One seed (seed=103) produced an ODE-fit failure (NaN parameter estimates), causing downstream PINN collapse ($R^2 = -28.9$). The 5-seed experiment (seeds 200–204, $R^2 = 0.945 \pm 0.011$) happened to avoid this failure mode. (2) Excluding the single ODE-fit failure, the 10-seed median (0.90) and the 5-seed mean (0.94) are consistent, confirming that Source 001 PINN performance is stable when the upstream ODE parameter estimation converges. (3) PINN is highly stable on curves where it works ($CV < 1\%$ for coating/surface curves).

Table 7 presents the complementary 5-seed $\times$ 18-curve experiment referenced in the Conclusion. PINN mean R^2 exceeds 0.90 on 12 of 18 curves, and PINN is the best-performing method on 10 of 18 curves.

4. Discussion

4.1. Why Pre-training Works—and Why Diversity Matters

The effectiveness of multi-mechanism pre-training can be understood through two complementary lenses. First, from a transfer learning perspective, the PINN learns a shared representation of fouling curve dynamics during pre-training [15]. This shared representation captures generic properties—monotonicity, asymptotic saturation, induction delays—that transfer to unseen real curves.

Second, the KS-large control experiment (Section 3.7) disentangles **data volume** from **mechanism diversity**. KS-small (200 trajectories) and KS-large (2,000 trajectories) both achieve near-zero pre-training loss ($\sim 1.5 \times 10^{-12}$), confirming that single-mechanism pre-training converges to functional-form memorization regardless of sample count. However, the downstream consequences differ: KS-large avoids catastrophic collapse—Source 001 R^2 improves from -1329 (KS-small) to $+0.62$ (KS-large)—showing that within-

Table 7: Multi-seed validation across all 18 curves ($N = 15$, 5 seeds, seeds 200–204). NC = Non-convergence.

Curve	PINN R^2	ODE-fit R^2	RF R^2	$R^2 > 0.9?$	Best
S001—Phosphoric acid	0.945 ± 0.011	0.986 ± 0.004	0.969 ± 0.023	Y	ODE
S002 CBA—Crude oil	0.258 ± 0.207	NC	0.539 ± 0.083	N	RF
S002 FED—Crude oil	0.642 ± 0.050	NC	0.648 ± 0.055	N	Tie*
S003 Run1—Unmodified	0.743 ± 0.005	NC	0.866 ± 0.044	N	RF
S003 Run2—Unmodified	0.545 ± 0.306	NC	0.753 ± 0.117	N	RF
S003 Run3—Unmodified	0.782 ± 0.068	NC	0.833 ± 0.094	N	RF
S003 Coating L1	0.995 ± 0.003	NC	0.976 ± 0.017	Y	PINN
S003 Coating L2	0.995 ± 0.004	NC	0.976 ± 0.011	Y	PINN
S003 Coating L3	0.997 ± 0.001	NC	0.992 ± 0.003	Y	PINN
S003 Coating Ref	0.993 ± 0.004	NC	0.990 ± 0.005	Y	PINN
S003 Fig.11 Ref	0.996 ± 0.001	NC	0.985 ± 0.008	Y	PINN
S003 Fig.11 Grit80	0.994 ± 0.002	NC	0.987 ± 0.003	Y	PINN
S003 Fig.11 Grit220	0.998 ± 0.000	NC	0.987 ± 0.005	Y	PINN
S003 Fig.11 DiaPro	0.990 ± 0.002	NC	0.981 ± 0.006	Y	PINN
S003 Fig.13 Flat	0.994 ± 0.004	NC	0.991 ± 0.001	Y	PINN
S003 Fig.13 Pattern1	0.964 ± 0.011	NC	0.979 ± 0.009	Y	RF
S003 Fig.13 Pattern2	0.949 ± 0.011	NC	0.974 ± 0.018	Y	RF
S003 Fig.13 Pattern3	0.824 ± 0.320	NC	0.761 ± 0.256	N	PINN

*Tie = statistically tied (overlapping $\pm 1\sigma$). PINN best on 10/18, RF on 7/18 (including 1 tie), ODE-fit on 1/18. Mean $R^2 > 0.90$ on 12/18 curves. Seed range differs from Table 6 (200–204 vs. 100–109); see Section 3.9 for discussion of seed variability.

mechanism data volume provides modest regularization. Multi-mechanism pre-training still outperforms KS-large by +0.29 on Source 001 ($R^2 = 0.91$ vs. 0.62), confirming that mechanism diversity provides gains that cannot be replicated by simply scaling up single-mechanism data volume. This finding parallels observations in multi-task learning [23], where training on diverse tasks produces representations that transfer better than those from any single task—even when the single-task training set is enlarged.

When is single-mechanism pre-training sufficient? KS-small outperforms multi-mechanism on three curves (Source 003 Run1–3, unmodified CaCO_3 , $R^2 = 0.83$ –0.92 vs. 0.55–0.78). However, KS-large—with equal data volume to multi-mechanism—underperforms KS-small on the same curves (Run1: 0.46 vs. 0.92), suggesting that for these specific high-noise curves, the smaller KS-small prior happens to provide a better inductive bias than either a larger single-mechanism prior or a diverse multi-mechanism prior. On 12 of 18 curves (all coating, surface treatment, and pattern surface curves), the pre-training strategy matters little ($R^2 > 0.99$ regardless of pre-training), indicating that when fine-tuning data is sufficiently informative, pre-training distribution is secondary. The practical implication: **multi-mechanism pre-training with sufficient volume (2,000 trajectories) is the most robust default, but the optimal pre-training strategy depends on both target curve characteristics and data volume—not diversity alone.**

4.2. When Does Physics Loss Help?

Adding physics loss during fine-tuning (PT+Phys vs. PT+NoPhys) provides minimal additional benefit. The pre-training already encodes the physics into the network weights; further regularization during fine-tuning is redundant. The exception is the NoPT+Phys ablation on OOD-2 curves, where physics alone sometimes matches the pre-trained model—this occurs when the induction model closely matches the true dynamics and the 15 sparse points are sufficient for parameter estimation.

4.3. Why Direct ODE Fitting Fails—and How PINN Helps

A central question is: *if we know the governing ODE, why not simply fit it to the sparse data?* The answer, revealed by our baseline experiment (Table 2), is that direct ODE fitting fails on 15 of 18 curves with only 15 sparse data points. The 3-parameter models are under-determined with sparse noisy data, causing the optimizer to converge to parameters that fit the training points but generalize poorly.

The PINN circumvents this failure mode not by having a better ODE, but by having a learned prior over plausible curve shapes. The 3,000 epochs of multi-mechanism pre-training encode a rich distribution of physically realistic $R_f(t)$ trajectories into the network weights. During fine-tuning, this prior acts as a regularizer: sparse data points steer the prediction toward the correct curve, while pre-trained weights prevent the model from collapsing to degenerate solutions.

4.4. Prior Mismatch on Clean Asymptotic Curves

Source 001 (phosphoric acid) illustrates a practical vulnerability of the pipeline rather than a fundamental limitation of the method: when the upstream ODE parameter estimation step fails (1 of 10 seeds,

or 10% of trials in our 10-seed experiment), the estimated parameters contain NaN values that propagate through condition encoding and cause the PINN fine-tuning to collapse. This is not a PINN failure—it is a `scipy.curve_fit` convergence failure on a specific unlucky draw of sparse points. When ODE estimation converges (9 of 10 seeds), PINN performance is stable (median $R^2 = 0.90$, IQR = $[0.84, 0.92]$). The fix for this vulnerability is straightforward: add a fallback in the pipeline that detects NaN parameters and re-initializes with default values or retries with different initial guesses. We leave this robustness improvement to future work. The 5-seed experiment (seeds 200–204, all converged) demonstrates expected performance when ODE estimation succeeds.

4.5. Limitations and Future Work

Current limitations: (1) **ODE model selection is not automated.** Step 1 of the fine-tuning pipeline requires a human to visually inspect the digitized curve and select the appropriate ODE model. The five candidate models include three Kern-Seaton variants with overlapping functional forms, making automatic selection via AIC/BIC unreliable when only $N = 5$ –30 sparse points are available (the likelihood estimates themselves are under-determined). Automating this step—either through a meta-classifier trained on the pre-training data or through development of a unified ODE that handles all curve types without manual selection—is a necessary prerequisite for fully hands-off industrial deployment. (2) Absolute extrapolation remains difficult—while PINN outperforms RF, the absolute R^2 on the extrapolated region remains negative. (3) Crude oil curves ($R^2 = 0.58$ – 0.62) reflect genuine industrial complexity and would benefit from additional input features unavailable in the digitized publications. (4) The current approach fine-tunes on a single curve; joint fine-tuning across multiple curves from the same plant would likely improve robustness. (5) The PINN contains 13,249 parameters while fine-tuning uses only 5–30 data points—a case of *benign overparameterization* where the pre-training phase provides strong implicit regularization [23]. The no-pre-training ablation (A3) confirms that the network architecture alone cannot provide sufficient regularization. (6) The induction model’s sigmoid smoothing parameter $\alpha = 10$ (yielding a transition width of ~ 0.2 time units) was insensitive in the range $[5, 20]$.

Future work: (1) Uncertainty quantification via Ensemble PINN or Monte Carlo Dropout. (2) Cross-device transfer learning within the same plant. (3) Multi-fidelity data fusion with CFD/HTRI simulations. (4) Online learning with concept drift detection. (5) Field validation with an industrial partner.

5. Conclusion

This study presented a multi-mechanism pre-training framework for physics-informed fouling resistance prediction. The key findings are:

1. **Multi-mechanism diversity provides gains beyond data volume.** A KS-large control (2,000 Kern-Seaton trajectories, matching multi-mechanism volume) confirms that mechanism diversity ($\Delta R^2 = +0.29$ on Source 001, $+0.68$ on Source 002 FED) cannot be replicated by increasing single-mechanism

data volume. Both data volume and mechanism diversity contribute independently to pre-training quality—volume prevents the most extreme over-fitting, while diversity provides the additional regularization needed for robust generalization.

2. **The framework achieves practical accuracy on complex fouling curves.** A 5-seed multi-seed validation across all 18 curves ($N = 15$) confirms that PINN mean R^2 exceeds 0.90 on 12 of 18 curves. A complementary 10-seed experiment on 5 representative curves reveals that Source 001 performance is stable (median $R^2 = 0.90$, IQR = $[0.84, 0.92]$) provided the upstream ODE parameter estimation converges—a condition met in 9 of 10 seeds. The remaining curves fall into categories where either direct ODE fitting is simpler and more accurate (Source 001) or where both PINN and RF provide partial solutions in industrially challenging regimes.
3. **Physics-constrained pre-training enables extrapolation.** The pre-trained PINN consistently outperforms RF in extrapolation tasks ($\Delta R^2 = +1.3$ to $+18.7$ at $N = 5$).
4. **The approach is feasible for industrial deployment with one remaining automation gap.** Pre-training requires only publicly available physics models; fine-tuning takes approximately 3 seconds per heat exchanger on an NVIDIA RTX 4060 GPU; and PINN-driven cleaning strategies reduce dimensionless costs relative to fixed-interval policies. The current bottleneck is the need for human judgment in selecting the appropriate ODE model for a given curve—automating this step, or developing a single unified ODE that adapts to all curve types, is the main engineering task remaining before fully hands-off deployment.

Data and Code Availability

All code for data generation, model training, and experiment reproduction will be made available at a public GitHub repository with a Zenodo DOI upon publication. Digitized real fouling curves and pre-trained model weights are included¹.

Acknowledgments

The author thanks [To be added after advisor confirmation] for manuscript review and feedback. This work was supported by independent research at the College of Chemical and Biological Engineering, Zhejiang University.

References

- [1] Müller-Steinhagen, H., Malayeri, M. R., & Watkinson, A. P. (2005). Fouling of heat exchangers—new approaches to solve an old problem. *Heat Transfer Engineering*, 26(1), 1–4.

¹ Anonymous repository URL will be provided during the review process.

- [2] Raissi, M., Perdikaris, P., & Karniadakis, G. E. (2019). Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational Physics*, 378, 686–707.
- [3] Karniadakis, G. E., Kevrekidis, I. G., Lu, L., Perdikaris, P., Wang, S., & Yang, L. (2021). Physics-informed machine learning. *Nature Reviews Physics*, 3(6), 422–440.
- [4] Kern, D. Q., & Seaton, R. E. (1959). A theoretical analysis of thermal surface fouling. *British Chemical Engineering*, 4(5), 258–262.
- [5] Ebert, W., & Panchal, C. B. (1997). Analysis of Exxon crude-oil-slip stream coking data. In *Fouling Mitigation of Industrial Heat-Exchange Equipment*.
- [6] Epstein, N. (1983). Thinking about heat transfer fouling: A 5×5 matrix. *Heat Transfer Engineering*, 4(1), 43–56.
- [7] Jradi, R., Fguiri, A., Marvillet, C., & Jeday, M. R. (2018). Tubular heat exchanger fouling in phosphoric acid concentration process. *Heat and Mass Transfer*, 54, 2489–2500.
- [8] Benyahia, F., et al. (2014). Study of the fouling deposit in the heat exchangers of Algiers refinery. *International Journal of Industrial Chemistry*, 5, 1–10.
- [9] Riihimäki, M., et al. (2011). Crystallization fouling on modified heat transfer surfaces. *Proceedings of International Conference on Heat Exchanger Fouling and Cleaning*, Crete, Greece.
- [10] Diaz-Bejarano, E., Coletti, F., & Macchietto, S. (2015). A new dynamic model of crude oil fouling deposits. *AIChE Journal*, 61(1), 233–250.
- [11] Al Ismaili, R., Lee, M. W., & Wilson, D. I. (2019). Optimisation of heat exchanger network cleaning schedules. *Computers & Chemical Engineering*, 121, 409–425.
- [12] Pogiatis, T., Ishiyama, E. M., Paterson, W. R., Vassiliadis, V. S., & Wilson, D. I. (2012). Identifying optimal cleaning cycles for heat exchangers subject to fouling and ageing. *Applied Energy*, 89(1), 60–66.
- [13] Goswami, S., Anitescu, C., Chakraborty, S., & Rabczuk, T. (2020). Transfer learning enhanced physics informed neural network for phase-field modeling of fracture. *Theoretical and Applied Fracture Mechanics*, 106, 102447.
- [14] Chakraborty, S. (2021). Transfer learning based multi-fidelity physics informed deep neural network. *Journal of Computational Physics*, 426, 109942.
- [15] Finn, C., Abbeel, P., & Levine, S. (2017). Model-agnostic meta-learning for fast adaptation of deep networks. *Proceedings of ICML*, 1126–1135.

- 383 [16] Cai, S., Wang, Z., Wang, S., Perdikaris, P., & Karniadakis, G. E. (2021). Physics-informed neural
384 networks for heat transfer problems. *Journal of Heat Transfer*, 143(6), 060801.
- 385 [17] Coletti, F., & Macchietto, S. (2011). A dynamic, distributed model of shell-and-tube heat exchangers
386 undergoing crude oil fouling. *Industrial & Engineering Chemistry Research*, 50(8), 4515–4533.
- 387 [18] Markowski, M., & Urbaniec, K. (2005). Optimal cleaning schedule for heat exchangers in a heat exchanger
388 network. *Applied Thermal Engineering*, 25(7), 1019–1032.
- 389 [19] Müller-Steinhagen, H., Malayeri, M. R., & Watkinson, A. P. (2011). Heat exchanger fouling: Mitigation
390 and cleaning strategies. *Heat Transfer Engineering*, 32(3-4), 189–196.
- 391 [20] Deshannavar, U. B., Rafeen, M. S., Ramasamy, M., & Subbarao, D. (2010). Crude oil fouling: A review.
392 *Journal of Applied Sciences*, 10(24), 3167–3174.
- 393 [21] Diaz-Bejarano, E., Coletti, F., & Macchietto, S. (2016). A model-based method for visualization, moni-
394 toring, and diagnosis of fouling in shell-and-tube heat exchangers. *Industrial & Engineering Chemistry
395 Research*, 55(35), 9372–9386.
- 396 [22] Wang, S., Teng, Y., & Perdikaris, P. (2021). Understanding and mitigating gradient flow pathologies in
397 physics-informed neural networks. *SIAM Journal on Scientific Computing*, 43(5), A3055–A3081.
- 398 [23] Neyshabur, B., Tomioka, R., & Srebro, N. (2015). In search of the real inductive bias: On the role of
399 implicit regularization in deep learning. *ICLR Workshop*.
- 400 [24] Sundaramoorthy, S., & Nirmala, G. (2018). Fouling factor estimation using artificial neural network.
401 *Chemical Product and Process Modeling*, 13(4), 20180018.
- 402 [25] Burnham, K. P., & Anderson, D. R. (2002). *Model Selection and Multimodel Inference: A Practical
403 Information-Theoretic Approach*. Springer, 2nd edition.